

# Computation with C++

Keith O'Hara

01/17/2018

# C & C++

- C is a low-level programming language above assembly
  - ▶ Examples: Gurobi, **R**, the Linux kernel, git, embedded systems.

# C & C++

- C is a low-level programming language above assembly
  - ▶ Examples: Gurobi, **R**, the Linux kernel, git, embedded systems.
  - ▶ Can inline assembly code (see OpenBLAS)

# C & C++

- C is a low-level programming language above assembly
  - ▶ Examples: Gurobi, **R**, the Linux kernel, git, embedded systems.
  - ▶ Can inline assembly code (see OpenBLAS)
  - ▶ TIOBE Index: C was the ‘fastest grower of 2017’ (!?!).

# C & C++

- C is a low-level programming language above assembly
  - ▶ Examples: Gurobi, **R**, the Linux kernel, git, embedded systems.
  - ▶ Can inline assembly code (see OpenBLAS)
  - ▶ TIOBE Index: C was the ‘fastest grower of 2017’ (!?!).
- C++ is sometimes described as ‘C with classes.’
  - ▶ But note:  $C \not\subset C++$ .

# C & C++

- C is a low-level programming language above assembly
  - ▶ Examples: Gurobi, **R**, the Linux kernel, git, embedded systems.
  - ▶ Can inline assembly code (see OpenBLAS)
  - ▶ TIOBE Index: C was the ‘fastest grower of 2017’ (!?!).
- C++ is sometimes described as ‘C with classes.’
  - ▶ But note:  $C \not\subset C++$ .
  - ▶ Memory management (constructors and destructors)

# C & C++

- C is a low-level programming language above assembly
  - ▶ Examples: Gurobi, **R**, the Linux kernel, git, embedded systems.
  - ▶ Can inline assembly code (see OpenBLAS)
  - ▶ TIOBE Index: C was the ‘fastest grower of 2017’ (!?!).
- C++ is sometimes described as ‘C with classes.’
  - ▶ But note:  $C \not\subset C++$ .
  - ▶ Memory management (constructors and destructors)
  - ▶ A lot of Machine Learning libraries are written in C++

# C & C++

- C is a low-level programming language above assembly
  - ▶ Examples: Gurobi, **R**, the Linux kernel, git, embedded systems.
  - ▶ Can inline assembly code (see OpenBLAS)
  - ▶ TIOBE Index: C was the ‘fastest grower of 2017’ (!?!).
- C++ is sometimes described as ‘C with classes.’
  - ▶ But note:  $C \not\subset C++$ .
  - ▶ Memory management (constructors and destructors)
  - ▶ A lot of Machine Learning libraries are written in C++
  - ▶ Modern C++: C++11 (and beyond)




# C & C++

- C is a low-level programming language above assembly
  - ▶ Examples: Gurobi, **R**, the Linux kernel, git, embedded systems.
  - ▶ Can inline assembly code (see OpenBLAS)
  - ▶ TIOBE Index: C was the ‘fastest grower of 2017’ (!?!).
- C++ is sometimes described as ‘C with classes.’
  - ▶ But note:  $C \not\subset C++$ .
  - ▶ Memory management (constructors and destructors)
  - ▶ A lot of Machine Learning libraries are written in C++
  - ▶ Modern C++: C++11 (and beyond)
- Both are static **type-defined** languages.  
double a = 8.0; compared to **R**’s ‘a = 8’

# C & C++

- C is a low-level programming language above assembly
  - ▶ Examples: Gurobi, **R**, the Linux kernel, git, embedded systems.
  - ▶ Can inline assembly code (see OpenBLAS)
  - ▶ TIOBE Index: C was the ‘fastest grower of 2017’ (!?!).
- C++ is sometimes described as ‘C with classes.’
  - ▶ But note: C  $\not\subset$  C++.
  - ▶ Memory management (constructors and destructors)
  - ▶ A lot of Machine Learning libraries are written in C++
  - ▶ Modern C++: C++11 (and beyond)
- Both are static **type-defined** languages.  
double a = 8.0; compared to **R**’s ‘a = 8’
- Since C++11:  
auto a = 8.0;

# Compiled vs Interpreted Languages



# Intro to C++

```
#include <iostream>

int main()
{
    double a = 3.0, b = 2.0;
    double c = a + b;
    std::cout << "a+b=" << c << std::endl;

    return 0;
}
```

- Put into a file called `my_prog.cpp`
- Need a compiler! `gcc` vs `clang` vs `icc`.
  - ▶ macOS users: get Xcode from the App Store
  - ▶ Windows users: Rtools includes the `gcc` chain
- Open a terminal and type:  
`g++ my_prog.cpp -o my_prog.out`
- Run with: `./my_prog.out`

# Intro to C++

- So what is the compiler actually doing?

```
g++ my_prog.cpp -o my_prog.out
```

- Compilation process:

- (1) **Preprocessor**: cleanup, evaluate preprocessor directives (e.g., `-DDO_IF_DEF`), and include files
- (2) **Compiler** will generate assembly code
- (3) **Assembler** will generate machine code
- (4) The **linker** will connect any required outside libraries to your program

Or

$C/C++ \rightsquigarrow \text{Cleaned } C/C++ \rightsquigarrow S \rightsquigarrow \text{Machine Code}$

- To stop at (2) use `-E`; for (3) use `-S`; and for (4) use `-c`. E.g.,  

```
g++ -S my_prog.cpp -o my_prog.S
```

# Example: Dijkstra

---

**Algorithm 1** Dijkstra

---

1: **procedure** DIJKSTRA( $s, \mathcal{X}, \mathcal{A}$ ) ▷ Solve for  $\mathbf{d}^*$   
2:   Define the pair  $(u_j, d_j) \in (\mathcal{X}, \mathbb{R})$ , where  $d_j := \text{dist}(u_s, u_j)$ .  
3:    $u_s \in \mathcal{X}$ ,  $\mathcal{V} = \mathcal{X}$ ;  $\mathbf{d}^* = \infty$ ,  $d_s^* = 0$  ▷ Initialization  
4:   **while**  $\mathcal{V} \neq \{\emptyset\}$  **do**  
5:     Choose  $u_i := \{u_j \in \mathcal{V} : d_j^* = \min_k \{d_k^*\}\}$ .  
6:     **if**  $d_j^* = \infty$  **then**  
7:       **break**;  
8:      $\mathcal{V} = \mathcal{V} \setminus \{u_i\}$  ▷ Remove  $u_i$   
9:     Define the adjacent network to  $u_i$  by  $N(u_i)$ .  
10:    For each  $v_j \in N(u_i) \cap \mathcal{V}$ , with arc weights  $\{a_{i,j}\}$ , calculate  
$$c_j = d_i^* + a_{i,j}$$
  
11:    **if**  $c_j < d_j^*$  **then**  
12:       $d_j^* = c_j$  ▷ Update  $d_j^*$   
13:    **return**  $\mathbf{d}^* = \{d_j^*\}_{j \in \mathcal{X}}$ . ▷ Solution

---

# Example: Dijkstra

- Recall the NYC roads example: approx 67k nodes, with 170k arcs



# Example: Dijkstra

- Recall the NYC roads example: approx 67k nodes, with 170k arcs
- Gurobi takes around 4 secs to solve for the min-dist path

## Example: Dijkstra

- Recall the NYC roads example: approx 67k nodes, with 170k arcs
- Gurobi takes around 4 secs to solve for the min-dist path
- SPR takes around 1/100 of a second

## Example: Dijkstra

- Recall the NYC roads example: approx 67k nodes, with 170k arcs
- Gurobi takes around 4 secs to solve for the min-dist path
- SPR takes around 1/100 of a second
- Why the difference?

## Example: Dijkstra

- Recall the NYC roads example: approx 67k nodes, with 170k arcs
- Gurobi takes around 4 secs to solve for the min-dist path
- SPR takes around 1/100 of a second
- Why the difference?
  - ▶ Efficient containers using the standard library: sparse data formats and efficiently ordered `std::set`.
  - ▶ C++ move instructions
  - ▶ Loop unrolling

# Loop Unrolling

# Loop Unrolling

- Loop ‘unrolling’ is touted as a major benefit of C/C++
- But takes some work to get right
- Parallelism on modern computers:
  - ▶ Single instruction, multiple data (SIMD) ; Streaming SIMD Extensions (SEE); Advanced Vector Extensions (AVX).
  - ▶ OpenMP and SIMD
- Compiler optimization flags:
  - O3 vs -Ofast; -march=native; -funroll-loops with gcc
- Generally you need to inspect the **assembly code!**

# Loop Unrolling: Example

- Consider the loop calculating:  $a(i) = a(i) + k \times b(i)$ ,  $i \in [n]$ .
- In C++ we could write this as:

```
void add_vecs(float* a, float* b, float k, int n)
{
    for (int i=0; i<n; i++) {
        a[i] += k*b[i];
    }
}
```

- Suppose  $n = 1\text{E}05$  and repeat 1000 times.
- (Optimized) **R** takes around 12 seconds.

# Loop Unrolling: Example

- Modify the code slightly with a hint to the compiler (`__restrict`) and an OpenMP pragma.

```
void add_vec(float* __restrict a, float* __restrict b,
             const float k, int n) {
    #pragma omp parallel for
    for (int i=0; i<n; i++) {
        a[i] += k*b[i];
    }
}
```



# Loop Unrolling: Example

- Modify the code slightly with a hint to the compiler (`__restrict`) and an OpenMP pragma.

```
void add_vec(float* __restrict a, float* __restrict b,
             const float k, int n) {
    #pragma omp parallel for
    for (int i=0; i<n; i++) {
        a[i] += k*b[i];
    }
}
```

- Let's inspect the assembly code! Compiled with

```
clang++-mp-5.0 -std=c++11 -Ofast -march=native -fopenmp \
simple_float_run.cpp -o simple_float.out
```

# Loop Unrolling: Example

- Modify the code slightly with a hint to the compiler (`__restrict`) and an OpenMP pragma.

```
void add_vec(float* __restrict a, float* __restrict b,
             const float k, int n) {
    #pragma omp parallel for
    for (int i=0; i<n; i++) {
        a[i] += k*b[i];
    }
}
```

- Let's inspect the assembly code! Compiled with

```
clang++-mp-5.0 -std=c++11 -Ofast -march=native -fopenmp \
simple_float_run.cpp -o simple_float.out
```

- Without OpenMP: 0.026 secs. With OpenMP: 0.01 secs.

# Loop Unrolling: Example

- Modify the code slightly with a hint to the compiler (`__restrict`) and an OpenMP pragma.

```
void add_vec(float* __restrict a, float* __restrict b,
             const float k, int n) {
    #pragma omp parallel for
    for (int i=0; i<n; i++) {
        a[i] += k*b[i];
    }
}
```

- Let's inspect the assembly code! Compiled with

```
clang++-mp-5.0 -std=c++11 -Ofast -march=native -fopenmp \
simple_float_run.cpp -o simple_float.out
```

- Without OpenMP: 0.026 secs. With OpenMP: 0.01 secs.
- tl;dr: **R** takes ( $>$ ) 1000 times longer than C++.

# Templates

# Template Programming

- **Templates** are an approach to generic programming in a type-defined language.
- Frequently used to avoid tedious function overloading.

Example: calculate the maximum of two numbers.

```
int max(int a, int b)
{
    int res = (a > b) ? a : b;
    return res;
}
```

```
// template version
template <typename T>
T max(T a, T b)
{
    return (a > b) ? a : b;
}
```

# Template Metaprogramming

- Second use: compile-time computation (vs run-time).

```
template <int n> struct Factorial {  
    static const int result = n * Factorial<n-1>::result;  
};  
  
// specialization in 0! case  
template <> struct Factorial<0> {  
    static const int result = 1;  
};  
  
int main()  
{  
    std::cout << Factorial<10>::result << std::endl;  
    return 0;  
}
```

- Pros and Cons?

# Template Metaprogramming

- New style with C++11:

```
constexpr
int
factorial(const int x) {
    return ( x == 0 ? 1 : x == 1 ? x : x*factorial(x-1));
}

int main()
{
    constexpr int x = 10;
    constexpr int res = factorial(x);

    return 0;
}
```

# Compile-time Computation

- Compilation with -O0:

```
_main:                                     ## @main
        .cfi_startproc
## BB#0:
        push    rbp
Lcfi0:
        .cfi_def_cfa_offset 16
Lcfi1:
        .cfi_offset rbp, -16
        mov     rbp, rsp
Lcfi2:
        .cfi_def_cfa_register rbp
        xor     eax, eax
        mov     dword ptr [rbp - 4], 0
        mov     dword ptr [rbp - 8], 10
        mov     dword ptr [rbp - 12], 3628800
        pop     rbp
        ret
        .cfi_endproc
```



# Compile-time Computation

- You might be thinking...

# Compile-time Computation

- You might be thinking...

That's a really silly example, Keith!

# Compile-time Computation

- You might be thinking...

That's a really silly example, Keith!

How many other examples could there be?

# Compile-time Computation

- You might be thinking...

That's a really silly example, Keith!

How many other examples could there be?

- Answer: A lot of functions possess recursive representations

# Compile-time Computation

- You might be thinking...

That's a really silly example, Keith!

How many other examples could there be?

- Answer: A lot of functions possess recursive representations
- Example:

$$\operatorname{erf}(x) = \frac{2}{\sqrt{\pi}} \int_0^x \exp(-t^2) dt$$

can be (well-) approximated using a continued fraction representation

$$\operatorname{erf}(x) = \frac{2x}{\sqrt{\pi}} \exp(-x^2) \cfrac{1}{1 - 2x^2 + \cfrac{4x^2}{3 - 2x^2 + \cfrac{8x^2}{5 - 2x^2 + \cfrac{12x^2}{7 - 2x^2 + \ddots}}}}$$

# Deeper Down the Rabbit Hole...

## OpenBLAS and Optimized GEMM

# OpenBLAS

- BLAS (Basic Linear Algebra Subprograms) and LAPACK (Linear Algebra PACKage)
- OpenBLAS is a fork of GotoBLAS (written by Kazushige Goto)
- **Hand-optimized assembly code for `gemm`**

# OpenBLAS

- BLAS (Basic Linear Algebra Subprograms) and LAPACK (Linear Algebra PACKage)
- OpenBLAS is a fork of GotoBLAS (written by Kazushige Goto)
- **Hand-optimized assembly code** for `gemm`
  - ▶ Optimization focused on use of L1 and L2 caches, TLB misses, and specialized architecture instructions



# OpenBLAS

- BLAS (Basic Linear Algebra Subprograms) and LAPACK (Linear Algebra PACKage)
- OpenBLAS is a fork of GotoBLAS (written by Kazushige Goto)
- **Hand-optimized assembly code** for `gemm`
  - ▶ Optimization focused on use of L1 and L2 caches, TLB misses, and specialized architecture instructions
- Get using:

```
git clone -b develop --single-branch \  
https://github.com/xianyi/OpenBLAS ~/Desktop/OBLAS
```

# OpenBLAS

- BLAS (Basic Linear Algebra Subprograms) and LAPACK (Linear Algebra PACKage)
- OpenBLAS is a fork of GotoBLAS (written by Kazushige Goto)
- **Hand-optimized assembly code** for `gemm`
  - ▶ Optimization focused on use of L1 and L2 caches, TLB misses, and specialized architecture instructions

- Get using:

```
git clone -b develop --single-branch \  
https://github.com/xianyi/OpenBLAS ~/Desktop/OBLAS
```

- Build using:

```
cd ~/Desktop/OBLAS  
make CC=gcc FC=gfortran  
sudo make PREFIX=/usr/local install
```

# OpenBLAS

- BLAS (Basic Linear Algebra Subprograms) and LAPACK (Linear Algebra PACKage)
- OpenBLAS is a fork of GotoBLAS (written by Kazushige Goto)
- **Hand-optimized assembly code** for `gemm`
  - ▶ Optimization focused on use of L1 and L2 caches, TLB misses, and specialized architecture instructions
- Get using:

```
git clone -b develop --single-branch \  
https://github.com/xianyi/OpenBLAS ~/Desktop/OBLAS
```
- Build using:

```
cd ~/Desktop/OBLAS  
make CC=gcc FC=gfortran  
sudo make PREFIX=/usr/local install
```
- Can easily link with **R** (replace Rblas or build **R** from source)

## Example: FMA Optimization

- One important aspect of hand-optimized assembly code relates to advanced processor instructions, such as fused multiply-add (FMA):

$$a \leftarrow b + c \times d$$

- In code:

## Example: FMA Optimization

- One important aspect of hand-optimized assembly code relates to advanced processor instructions, such as fused multiply-add (FMA):

$$a \leftarrow b + c \times d$$

- In code:

```
double add_m3(double x1, double x2, double x3)
{
    return x1 + x2*x3;
}
```

## Example: FMA Optimization

- One important aspect of hand-optimized assembly code relates to advanced processor instructions, such as fused multiply-add (FMA):

$$a \leftarrow b + c \times d$$

- In code:

```
double add_m3(double x1, double x2, double x3)
{
    return x1 + x2*x3;
}
```

- Using `-O3`

```
mulsd    xmm1, xmm2
addsd    xmm0, xmm1
pop      rbp
ret
```

## Example: FMA Optimization

- One important aspect of hand-optimized assembly code relates to advanced processor instructions, such as fused multiply-add (FMA):

$$a \leftarrow b + c \times d$$

- In code:

```
double add_m3(double x1, double x2, double x3)
{
    return x1 + x2*x3;
}
```

- Using `-Ofast -march=native`

```
    vfmadd231sd    xmm0, xmm2, xmm1
    pop            rbp
    ret
```

# Linear Algebra with C++: Armadillo



# Armadillo

- Armadillo is a **templated C++** linear algebra library
- Written by Conrad Sanderson at NICTA/Data61
- **Template metaprogramming** & **static polymorphism**
- Syntax similar to Matlab's

Armadillo	Matlab
$C = A.t() * B$	$C = A' * B$
$C = A \% B$	$C = A .* B$
$C = \text{solve}(A, B)$	$C = A \backslash B$
$\text{int } n = A.n\_rows$	$n = \text{size}(A, 1)$

- Link against **OpenBLAS** or some other system BLAS/LAPACK

# Armadillo + OpenBLAS

- Consider

$$Z = A \times B \times C \times D$$

where the dimensions are

$$A = 1000 \times 800, \quad B = 800 \times 600$$

$$C = 600 \times 400, \quad D = 400 \times 200$$

# Armadillo + OpenBLAS

- Consider

$$Z = A \times B \times C \times D$$

where the dimensions are

$$A = 1000 \times 800, \quad B = 800 \times 600$$

$$C = 600 \times 400, \quad D = 400 \times 200$$

- How do we compute this efficiently?

# Armadillo + OpenBLAS

- Consider

$$Z = A \times B \times C \times D$$

where the dimensions are

$$A = 1000 \times 800, \quad B = 800 \times 600$$

$$C = 600 \times 400, \quad D = 400 \times 200$$

- How do we compute this efficiently?
- Do this 1000 times.

Run Time (sec.)		
<b>R</b>	<b>R &amp; OpenBLAS</b>	<b>C++ &amp; OpenBLAS</b>
475.208	18.01	3.96